

```
In [2]: import numpy as np
```

```
In [3]: NUM_FEATS = 4
        NUM_STATES = 4
        NUM_ACTIONS = 2
        GAMMA = 0.9
```

```
In [4]: # Expected rewards for each state action pair
        R = np.array([-40, -80, -160, -380, -100, -100, -100, -100])
```

```
In [5]: c = np.ones(NUM_STATES)
        PHI = np.eye(NUM_FEATS)
        P_0 = np.array(
            [
                [0.5, 0.5, 0, 0],
                [0, 0.5, 0.5, 0],
                [0, 0, 0.5, 0.5],
                [0.5, 0, 0, 0.5]
            ]
        )
        P_1 = np.concatenate((np.ones((4,1)), np.zeros((4, 3))), axis=1)
        A = np.concatenate((np.eye(4), np.eye(4))) - GAMMA * np.concatenate((P_0, P_1))
        c, PHI, A
```

```
Out[5]: (array([1., 1., 1., 1.]),
        array([[1., 0., 0., 0.],
               [0., 1., 0., 0.],
               [0., 0., 1., 0.],
               [0., 0., 0., 1.]]),
        array([[ 0.55, -0.45,  0.,  0. ],
               [ 0.,  0.55, -0.45,  0. ],
               [ 0.,  0.,  0.55, -0.45],
               [-0.45,  0.,  0.,  0.55],
               [ 0.1 ,  0.,  0.,  0. ],
               [-0.9 ,  1.,  0.,  0. ],
               [-0.9 ,  0.,  1.,  0. ],
               [-0.9 ,  0.,  0.,  1. ]]))
```

```
In [6]: PHI_HAT = np.eye(8)
        R_TILDE = np.concatenate((np.diag(R[0:4]), np.diag(R[4:])), axis=0)
        PHI_HAT, R_TILDE
```

```
Out[6]: (array([[1., 0., 0., 0., 0., 0., 0., 0.],
               [0., 1., 0., 0., 0., 0., 0., 0.],
               [0., 0., 1., 0., 0., 0., 0., 0.],
               [0., 0., 0., 1., 0., 0., 0., 0.],
               [0., 0., 0., 0., 1., 0., 0., 0.],
               [0., 0., 0., 0., 0., 1., 0., 0.],
               [0., 0., 0., 0., 0., 0., 1., 0.],
               [0., 0., 0., 0., 0., 0., 0., 1.]]),
        array([[ -40,  0,  0,  0],
               [  0, -80,  0,  0],
               [  0,  0, -160,  0],
               [  0,  0,  0, -380],
               [-100,  0,  0,  0],
```

```
[ 0, -100, 0, 0],
[ 0, 0, -100, 0],
[ 0, 0, 0, -100]]))
```

In [7]:

```
import gurobipy as gp
from gurobipy import GRB

m = gp.Model('lp')

w_z = m.addMVar(8, lb=float('-inf'))
c = c.T @ PHI
m.setMObjective(None, c=c, constant=0.0, xc=w_z[0:4], sense=GRB.MINIMIZE)

m.addConstr(A @ PHI @ w_z[0:4] >= PHI_HAT @ R_TILDE @ w_z[4:])
m.addConstr(sum(w_z[4:]) == 1)
m.addConstrs(w_z[i] >= 0 for i in range(4, 8))

m.optimize()
```

Academic license - for non-commercial use only - expires 2021-05-28
 Using license file C:\Users\Brendan\gurobi.lic
 Gurobi Optimizer version 9.1.1 build v9.1.1rc0 (win64)
 Thread count: 4 physical cores, 8 logical processors, using up to 8 threads
 Optimize a model with 13 rows, 8 columns and 31 nonzeros
 Model fingerprint: 0xa4996764
 Coefficient statistics:
 Matrix range [1e-01, 4e+02]
 Objective range [1e+00, 1e+00]
 Bounds range [0e+00, 0e+00]
 RHS range [1e+00, 1e+00]
 Presolve removed 8 rows and 0 columns
 Presolve time: 0.04s
 Presolved: 5 rows, 8 columns, 30 nonzeros

Iteration	Objective	Primal Inf.	Dual Inf.	Time
0	-3.2388430e+03	1.774277e+01	0.000000e+00	0s
4	-9.4815014e+02	0.000000e+00	0.000000e+00	0s

Solved in 4 iterations and 0.07 seconds
 Optimal objective -9.481501404e+02

In [8]:

```
w_opt = np.array([v.x for v in m.getVars()])[0:4]
```

In [9]:

```
PHI @ w_opt
```

Out[9]: array([-236.00208435, -267.46902893, -229.00943 , -215.66959708])